

traffic-accidents-analysis-and-ml

March 27, 2025

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
pd.set_option("display.max_columns",None)
pd.set_option('display.float_format', '{:.2f}'.format)
```

```
[2]: from sklearn.linear_model import Ridge, Lasso, ElasticNet , LinearRegression
from sklearn.model_selection import RandomizedSearchCV , cross_validate ,_
    ↪train_test_split
from xgboost import XGBRegressor
from sklearn.metrics import r2_score , mean_squared_error
from keras.models import Sequential
from keras.layers import Dense
import tensorflow as tf
from catboost import CatBoostRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.linear_model import BayesianRidge
from sklearn.preprocessing import MinMaxScaler
from sklearn.ensemble import VotingRegressor
```

```
[3]: data = pd.read_csv("traffic_accidents_dict_new.csv")
df = data.copy()
df.head()
```

```
[3]:   accidents  traffic_fine_amount  traffic_density  traffic_lights \
0          20                  4.37          2.30      753.00
1          11                  9.56          3.28       5.45
2          19                  7.59          2.10       6.70
3          23                  6.39          4.92       9.41
4          23                  2.40          1.96       7.39
```

```

pavement_quality  urban_area  average_speed  rain_intensity  vehicle_count \
0                 0.77        1             321.59          1.19           290.86
1                 4.05        1             478.62          6.30           931.81
2                345.00        0             364.48          2.86           830.09
3                 4.73        0             20.92          2.11           813.16
4                 1.71        1             37.38          1.70           1.47

      time_of_day           Variable \
0            160.43       accidents
1            8.91    traffic_fine_amount
2            5.57    traffic_density
3           131.45    traffic_lights
4            6.96    pavement_quality

                           Description
0 Number of recorded accidents, as a positive in...
1 Traffic fine amount, expressed in thousands of...
2 Traffic density index, scale from 0 (low) to 1...
3 Proportion of traffic lights in the area (0 to...
4 Pavement quality, scale from 0 (very poor) to ...

```

```
[4]: def check_df(data,head=False):
    print("----- SHAPE -----")
    print(data.shape)
    print("----- INFO -----")
    print(data.info())
    print("----- ISNA -----")
    print(data.isnull().sum().sort_values(ascending=False))
    print("----- NUNIQUE -----")
    print(data.nunique().sort_values())
    print("----- DUPLICATED -----")
    print(data.duplicated().sum())
    print("----- DESCRIBE -----")
    display(data.describe().T)
    if head:
        print("----- HEAD -----")
        display(data.head())
        print("----- TAIL -----")
        display(data.tail())

```

```
[5]: def col_types(data):
    cat_cols = [col for col in data.columns if data[col].dtypes=="O"]
    num_cols = [col for col in data.columns if data[col].dtypes!="O"]
    print(f"OBSERVATION : {data.shape[1]}")
    print(f"CAT COLS : {len(cat_cols)}")
    print(f"NUM COLS : {len(num_cols)}")
    return cat_cols , num_cols

```

```
[6]: def box_plot(df,col):
    sns.boxplot(data=df,x=col,palette="inferno")
    plt.title(f'{col.upper()} boxplot')
    plt.tight_layout()
    plt.show()

[7]: def outlier_thresholds(dataframe, col_name, q1=0.25, q3=0.75):
    quantile1 = dataframe[col_name].quantile(q1)
    quantile3 = dataframe[col_name].quantile(q3)
    interquantile_range = quantile3 - quantile1
    up_limit = quantile3 + 1.5 * interquantile_range
    low_limit = quantile1 - 1.5 * interquantile_range
    return low_limit, up_limit

[8]: def check_outlier(dataframe, col_name):
    low_limit, up_limit = outlier_thresholds(dataframe, col_name)
    if dataframe[(dataframe[col_name] > up_limit) | (dataframe[col_name] <
    ↪low_limit)].any(axis=None):
        return True
    else:
        return False

[9]: def replace_with_thresholds(dataframe, variable):
    low_limit, up_limit = outlier_thresholds(dataframe, variable)
    dataframe.loc[(dataframe[variable] < low_limit), variable] = low_limit
    dataframe.loc[(dataframe[variable] > up_limit), variable] = up_limit

[10]: def cat_plot(df, col):
    plt.figure(figsize=(14,6))
    order = df[col].value_counts().index
    ax = sns.countplot(x=col, data=df, palette="viridis", order=order)
    total = len(df[col])
    for p in ax.patches:
        height = p.get_height()
        percentage = (height / total) * 100
        ax.text(p.get_x() + p.get_width() / 2, height + 1, f'{percentage:.2f}%',
                ha="center", va="bottom", fontsize=10)
    plt.title(col.upper())
    plt.ylabel(col.upper())
    plt.xticks(rotation=45)
    plt.xlabel("COUNT")
    plt.tight_layout()
    plt.show()

    ↪
    ↪print("-----")
    ↪end="\n\n")
```

```
[11]: def hist_plot(data, col):
    plt.figure(figsize=(10,6))
    sns.histplot(x=col,data=df,kde=True,bins=25,color="g")
    plt.title(col.upper())
    plt.ylabel("COUNT")
    plt.xlabel(col.upper())
    plt.tight_layout()
    plt.show()

    □
    ↵print("-----")
```



```
[12]: def base_models(X,y,scoring="neg_mean_squared_error"):
    print("Base Models")
    regressions = [ ('Ridge',Ridge(alpha=1)),
                    ('Lasso',Lasso(alpha=1)),
                    ('ElasticNet',ElasticNet(alpha=1,l1_ratio=0.5)),
                    ("Linear",LinearRegression()),
                    ("Decision Tree",DecisionTreeRegressor()),
                    ("Random Forest",RandomForestRegressor()),
                    ('Gradient Boosting', GradientBoostingRegressor()),
                    ("XGB",XGBRegressor(objective='reg:squarederror')),
                    ('KNN', KNeighborsRegressor()),
                    ('MLP', MLPRegressor()),
                    ('Bayesian Ridge', BayesianRidge()),
                    ("CatBoost",CatBoostRegressor(verbose=0))]

    for name,regression in regressions:
        cv_results = cross_validate(regression,X,y,cv=10,scoring=scoring)
        print(f"{scoring}: {round(cv_results['test_score'].mean(),4)} {name}")
```



```
[13]: def hyperparameter_optimization_randomized(X, y, cv=5, scoring="r2", □
    ↵n_iter=100):
    print("HYPERPARAMETER OPTIMIZATION WITH RANDOMIZEDSEARCH")
    best_models = {}

    regressors = [
        ('GradientBoosting', GradientBoostingRegressor(), gb_params),
        ('CatBoost', CatBoostRegressor(verbose=0), cat_params),
        ('XGBoost', XGBRegressor(objective='reg:squarederror'), xgb_params)
    ]

    for name, regressor, params in regressors:
        print(f"----- {name} -----")

        cv_results = cross_validate(regressor, X, y, cv=cv, scoring=scoring)
```

```

        print(f"{scoring} (Before): {round(cv_results['test_score'].mean(), 4)}")

    rs_best = RandomizedSearchCV(
        regressor,
        params,
        cv=cv,
        n_jobs=-1,
        verbose=False,
        n_iter=n_iter,
        random_state=42
    ).fit(X, y)

    final_model = regressor.set_params(**rs_best.best_params_)

    cv_results = cross_validate(final_model, X, y, cv=cv, scoring=scoring)
    print(f"{scoring} (After): {round(cv_results['test_score'].mean(), 4)}")
    print(f"name best params: {rs_best.best_params_}", end="\n\n")

    best_models[name] = final_model

    return best_models

```

```
[14]: def voting_regressor(best_models, X, y, cv=5):
    print("Voting Regressor")

    voting_reg = VotingRegressor(estimators=[
        ("XGB", best_models['XGBoost']),
        ("GradientBoosting", best_models['GradientBoosting']),
        ("CatBoost", best_models["CatBoost"])
    ]).fit(X, y)

    cv_results = cross_validate(voting_reg, X, y, cv=cv, scoring=["neg_mean_squared_error", "r2"])

    print(f"Neg MSE: {-cv_results['test_neg_mean_squared_error'].mean()}")
    print(f"R^2: {cv_results['test_r2'].mean()}")

    return voting_reg
```

```
[15]: def plot_importance_ensemble(voting_reg, features):
    num=len(features)
```

```

feature_imp_df = pd.DataFrame(0, index=features.columns, columns=["Value"])

for model in voting_reg.estimators_:
    if hasattr(model, 'feature_importances_'):

        feature_imp_df["Value"] += model.feature_importances_ /100

feature_imp_df["Feature"] = features.columns
feature_imp_df = feature_imp_df.sort_values(by="Value", ascending=False)

plt.figure(figsize=(12, 7))
sns.set(font_scale=1)
ax = sns.barplot(x="Value", y="Feature", data=feature_imp_df[0:num])

for p in ax.patches:
    ax.annotate(format(p.get_width(), '.4f'),
                (p.get_width(), p.get_y() + p.get_height() / 2),
                ha='left', va='center', fontsize=8)

plt.title("Feature Importance - Voting Regressor")
plt.tight_layout()
plt.show()

```

[16]: check_df(df, True)

```

----- SHAPE -----
(8756, 12)
----- INFO -----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8756 entries, 0 to 8755
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   accidents        8756 non-null   int64  
 1   traffic_fine_amount 8756 non-null   float64 
 2   traffic_density   8756 non-null   float64 
 3   traffic_lights   8756 non-null   float64 
 4   pavement_quality  8756 non-null   float64 
 5   urban_area        8756 non-null   int64  
 6   average_speed     8756 non-null   float64 
 7   rain_intensity    8756 non-null   float64 
 8   vehicle_count     8756 non-null   float64 
 9   time_of_day       8756 non-null   float64 
 10  Variable          8756 non-null   object  
 11  Description        8756 non-null   object 

```

```
dtypes: float64(8), int64(2), object(2)
```

```
memory usage: 821.0+ KB
```

```
None
```

```
----- ISNA -----
```

```
accidents          0  
traffic_fine_amount 0  
traffic_density    0  
traffic_lights     0  
pavement_quality   0  
urban_area         0  
average_speed      0  
rain_intensity     0  
vehicle_count      0  
time_of_day        0  
Variable           0  
Description         0  
dtype: int64
```

```
----- NUNIQUE -----
```

```
urban_area         2  
Variable           11  
Description         12  
accidents          31  
traffic_lights     5600  
rain_intensity     7416  
pavement_quality   7892  
traffic_fine_amount 8326  
traffic_density     8330  
time_of_day        8586  
average_speed       8675  
vehicle_count       8717  
dtype: int64
```

```
----- DUPLICATED -----
```

```
0
```

```
----- DESCRIBE -----
```

	count	mean	std	min	25%	50%	75%	max
accidents	8756.00	20.63	5.23	5.00	17.00	21.00	24.00	35.00
traffic_fine_amount	8756.00	5.45	2.60	1.00	3.20	5.43	7.68	10.00
traffic_density	8756.00	14.25	75.97	0.24	3.32	5.56	7.89	996.00
traffic_lights	8756.00	93.29	224.17	0.00	3.78	6.50	9.23	999.00
pavement_quality	8756.00	22.44	112.19	0.00	2.10	3.28	4.43	994.00
urban_area	8756.00	0.69	0.46	0.00	0.00	1.00	1.00	1.00
average_speed	8756.00	214.64	168.71	0.97	12.51	223.13	360.94	932.00
rain_intensity	8756.00	33.86	140.05	0.00	1.67	2.36	3.70	999.00
vehicle_count	8756.00	453.23	313.88	1.03	169.46	453.98	729.15	999.62
time_of_day	8756.00	83.65	97.82	0.12	6.51	12.18	174.78	999.00

```
----- HEAD -----
```

	accidents	traffic_fine_amount	traffic_density	traffic_lights	\
0	20	4.37	2.30	753.00	
1	11	9.56	3.28	5.45	
2	19	7.59	2.10	6.70	
3	23	6.39	4.92	9.41	
4	23	2.40	1.96	7.39	

	pavement_quality	urban_area	average_speed	rain_intensity	vehicle_count	\
0	0.77	1	321.59	1.19	290.86	
1	4.05	1	478.62	6.30	931.81	
2	345.00	0	364.48	2.86	830.09	
3	4.73	0	20.92	2.11	813.16	
4	1.71	1	37.38	1.70	1.47	

	time_of_day	Variable	\
0	160.43	accidents	
1	8.91	traffic_fine_amount	
2	5.57	traffic_density	
3	131.45	traffic_lights	
4	6.96	pavement_quality	

	Description				
0	Number of recorded accidents, as a positive in...				
1	Traffic fine amount, expressed in thousands of...				
2	Traffic density index, scale from 0 (low) to 1...				
3	Proportion of traffic lights in the area (0 to...				
4	Pavement quality, scale from 0 (very poor) to ...				

----- TAIL -----

	accidents	traffic_fine_amount	traffic_density	traffic_lights	\
8751	27	2.77	6.94	6.60	
8752	18	3.26	1.73	549.00	
8753	31	2.56	8.53	2.78	
8754	10	9.62	1.40	2.72	
8755	12	8.69	4.20	5.39	

	pavement_quality	urban_area	average_speed	rain_intensity	\
8751	3.36	1	302.00	1.98	
8752	2.85	1	6.30	8.27	
8753	3.28	1	479.93	2.82	
8754	5.78	1	280.78	6.40	
8755	2.38	0	234.63	1.12	

	vehicle_count	time_of_day	Variable	Description
8751	258.92	11.15	0	0
8752	296.40	17.04	0	0
8753	458.21	192.77	0	0
8754	147.66	3.64	0	0

```
8755      483.98      1.08      0      0
```

```
[17]: cat_cols , num_cols = col_types(df)
```

```
OBSERVISION : 12
```

```
CAT COLS : 2
```

```
NUM COLS : 10
```

```
[18]: df[cat_cols].head()
```

```
[18]:           Variable                      Description
0       accidents    Number of recorded accidents, as a positive in...
1  traffic_fine_amount  Traffic fine amount, expressed in thousands of...
2       traffic_density  Traffic density index, scale from 0 (low) to 1...
3       traffic_lights  Proportion of traffic lights in the area (0 to...
4     pavement_quality  Pavement quality, scale from 0 (very poor) to ...
```

```
[19]: df[num_cols].head()
```

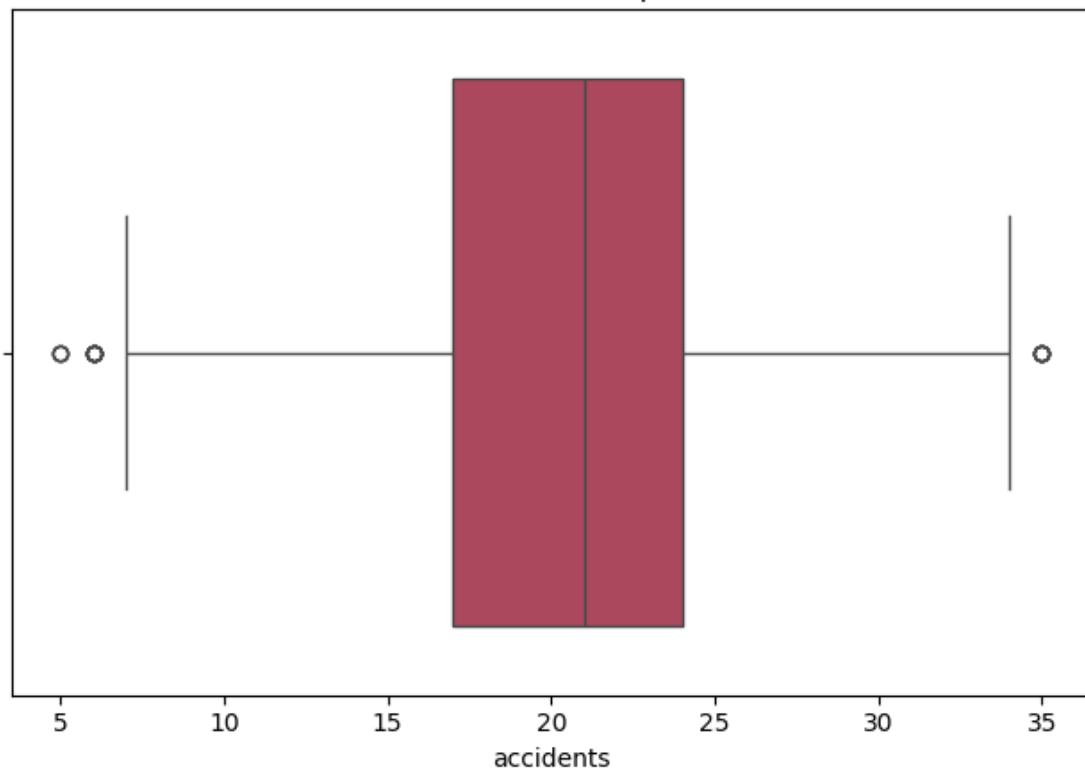
```
[19]:   accidents  traffic_fine_amount  traffic_density  traffic_lights \
0          20            4.37            2.30         753.00
1          11            9.56            3.28          5.45
2          19            7.59            2.10          6.70
3          23            6.39            4.92          9.41
4          23            2.40            1.96          7.39

  pavement_quality  urban_area  average_speed  rain_intensity  vehicle_count \
0            0.77            1        321.59          1.19        290.86
1            4.05            1        478.62          6.30        931.81
2           345.00            0        364.48          2.86        830.09
3            4.73            0        20.92          2.11        813.16
4            1.71            1        37.38          1.70          1.47

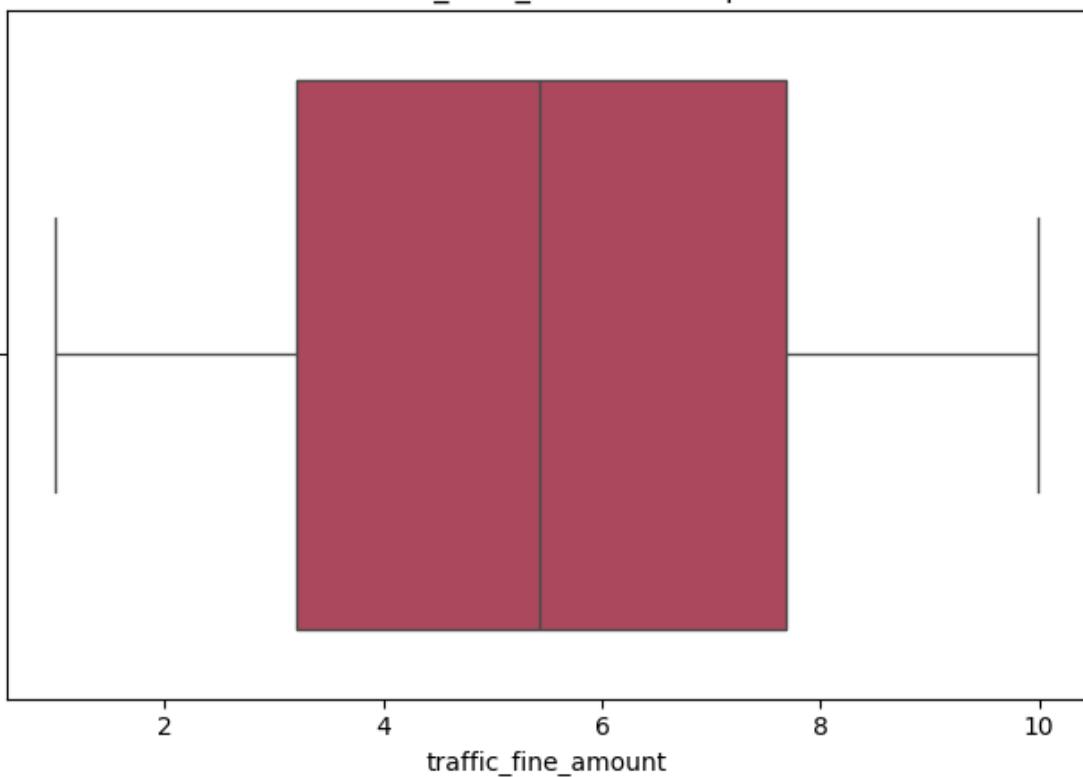
  time_of_day
0        160.43
1         8.91
2         5.57
3       131.45
4         6.96
```

```
[20]: for col in num_cols:
    box_plot(df,col)
```

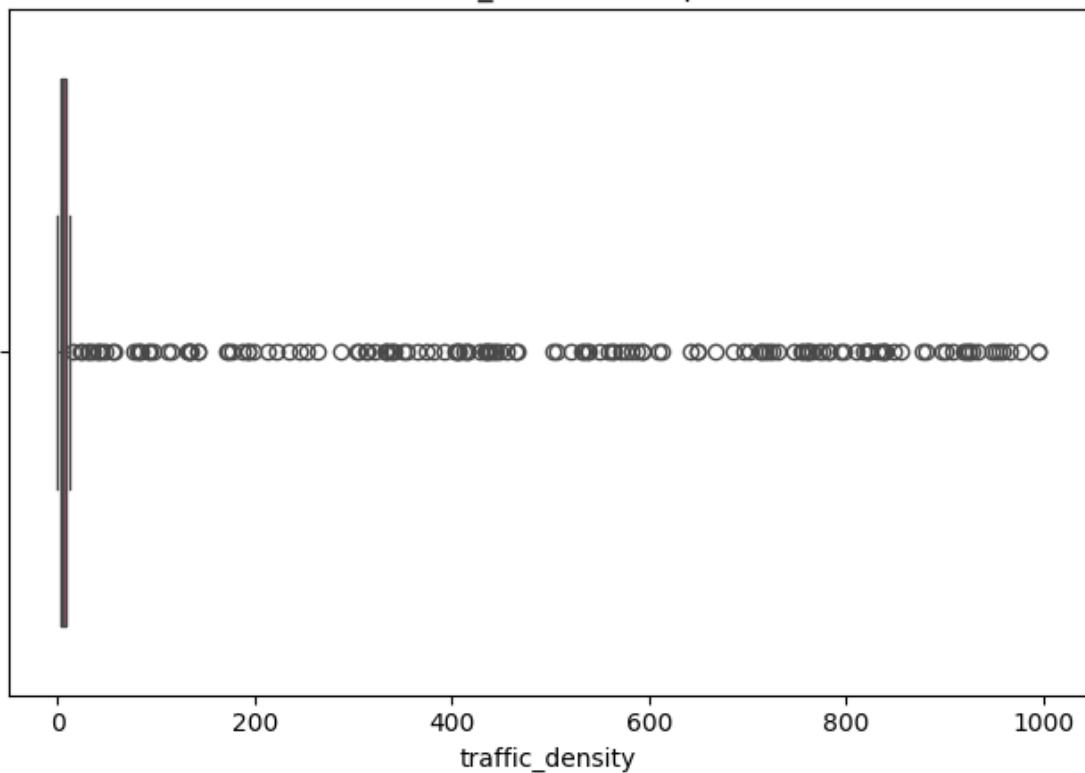
ACCIDENTS boxplot



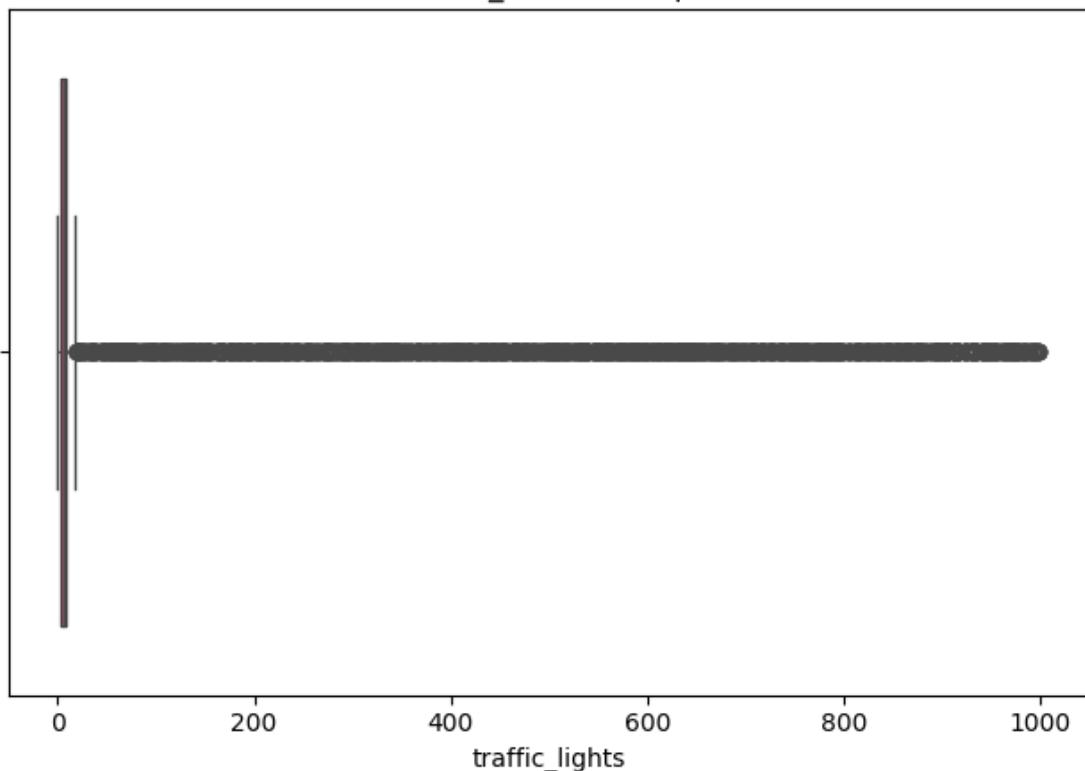
TRAFFIC_FINE_AMOUNT boxplot



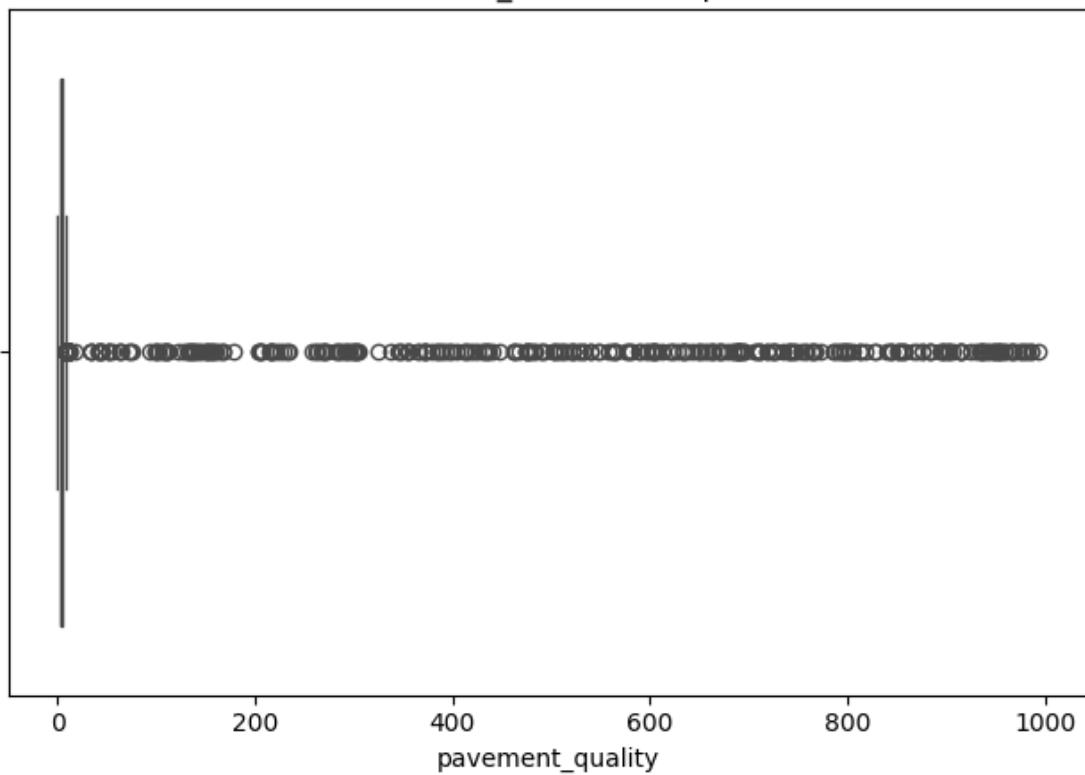
TRAFFIC_DENSITY boxplot



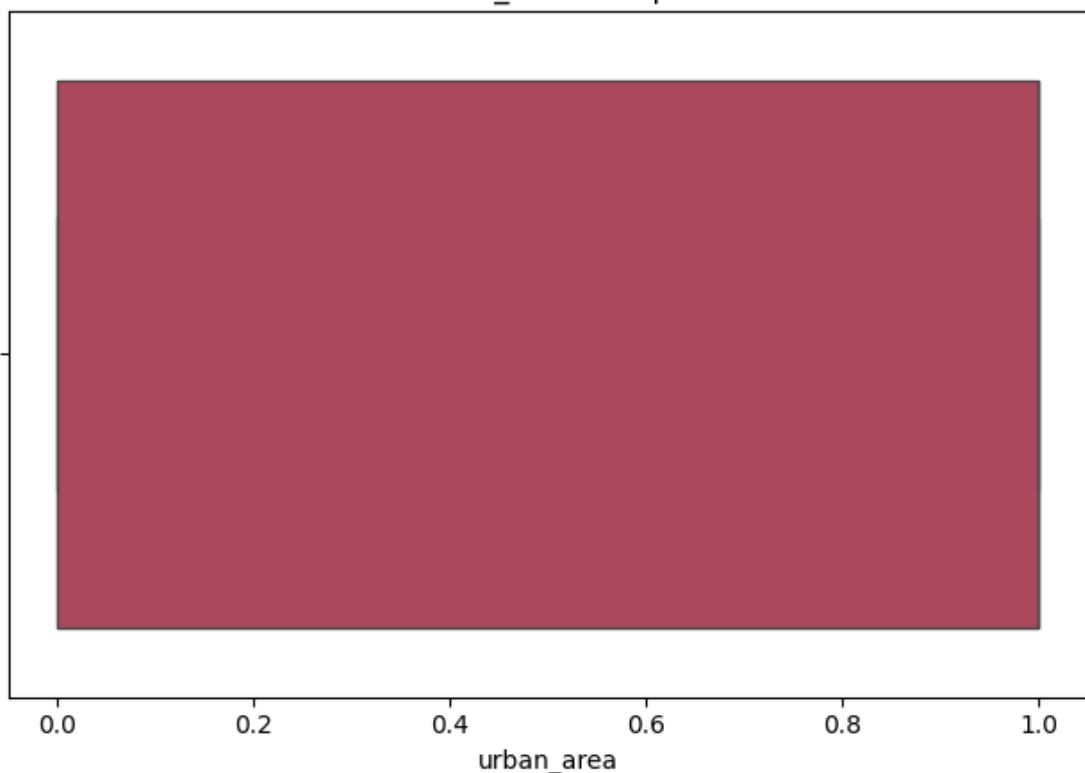
TRAFFIC_LIGHTS boxplot



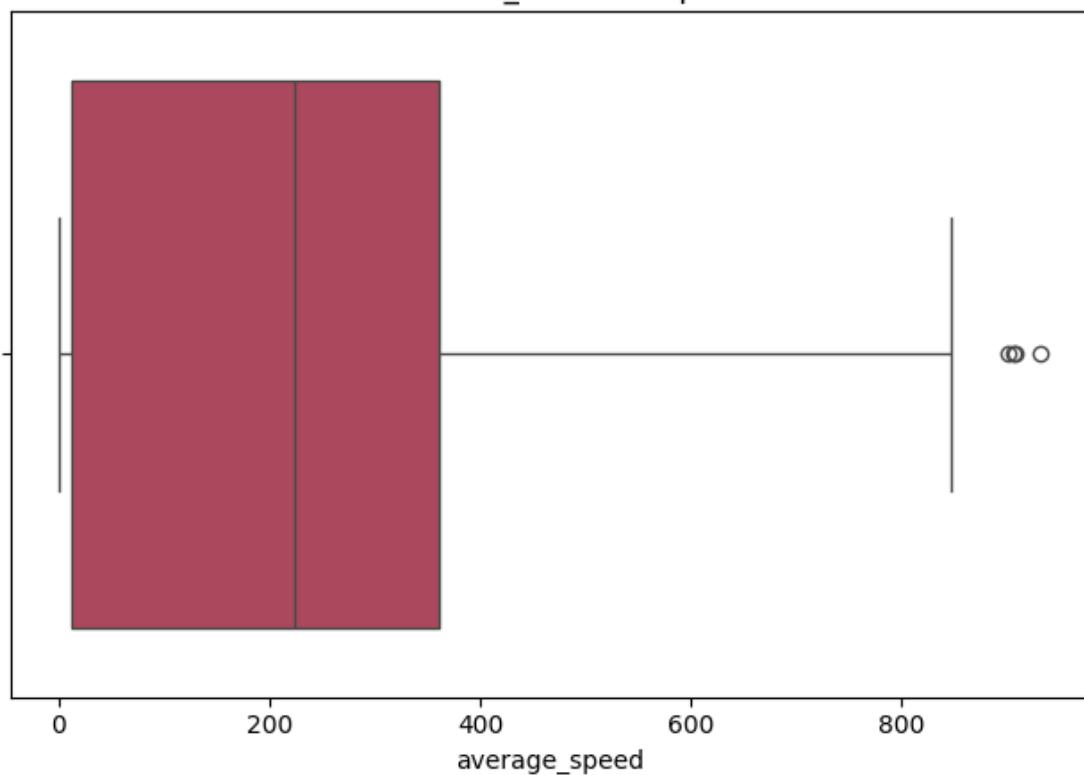
PAVEMENT_QUALITY boxplot



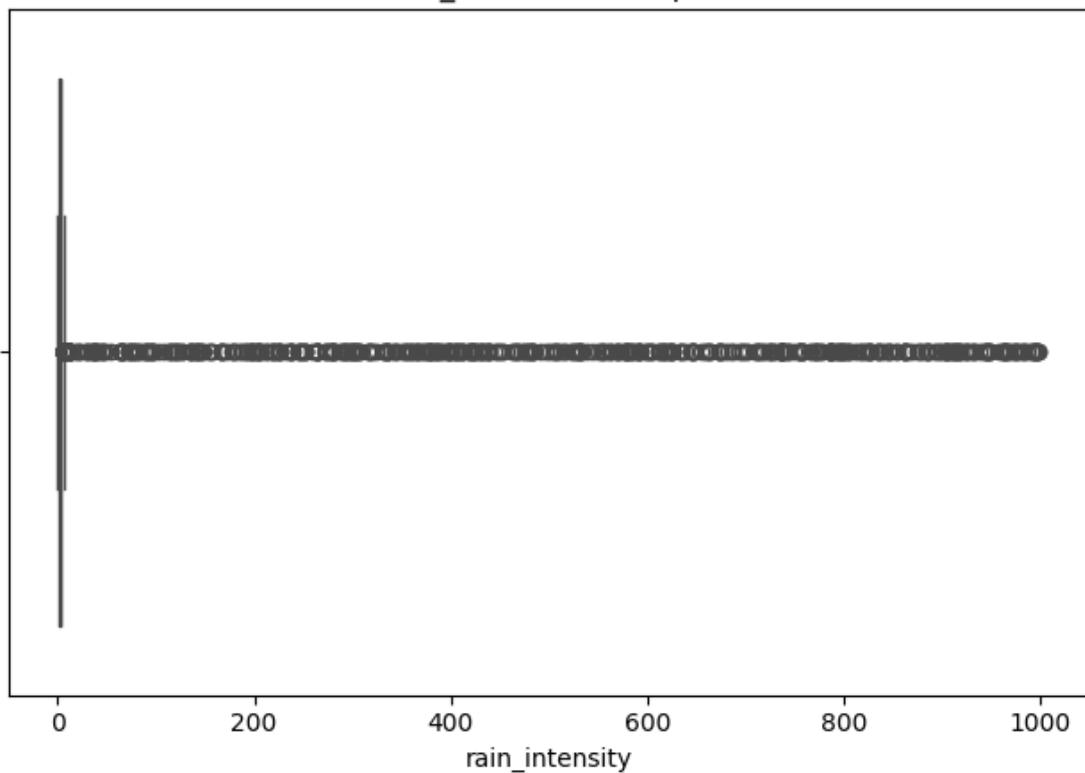
URBAN_AREA boxplot



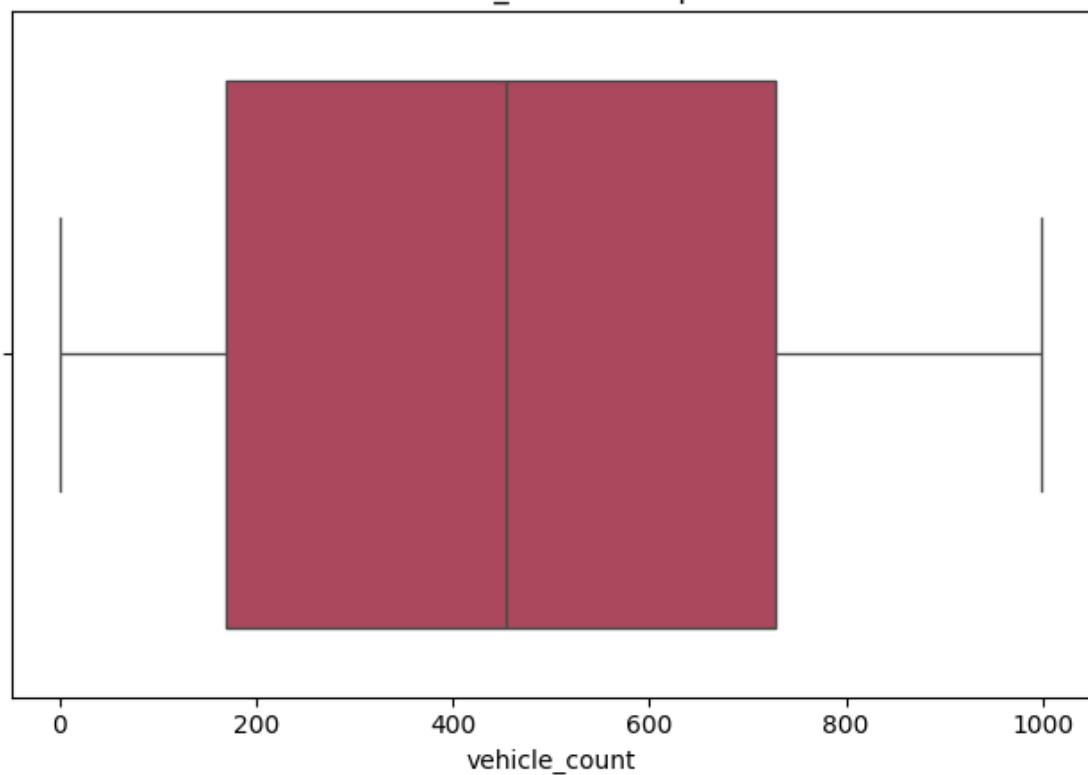
AVERAGE_SPEED boxplot

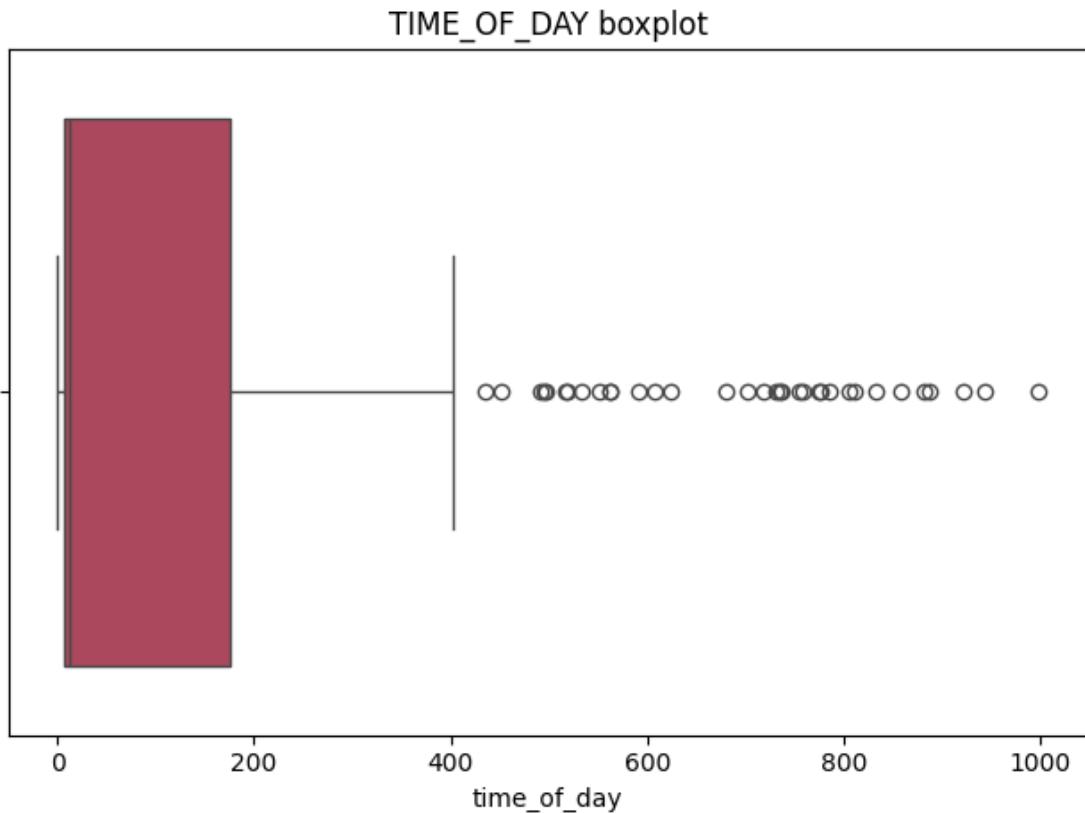


RAIN_INTENSITY boxplot



VEHICLE_COUNT boxplot





```
[21]: for col in num_cols:
    print(col.upper(), " ", check_outlier(df,col))
```

```
ACCIDENTS      True
TRAFFIC_FINE_AMOUNT  False
TRAFFIC_DENSITY      True
TRAFFIC_LIGHTS      True
PAVEMENT_QUALITY    True
URBAN_AREA        False
AVERAGE_SPEED      True
RAIN_INTENSITY      True
VEHICLE_COUNT      False
TIME_OF_DAY        True
```

```
[22]: for col in num_cols:
    replace_with_thresholds(df, col)
```

```
[23]: for col in num_cols:
    print(col.upper(), " ", check_outlier(df,col))
```

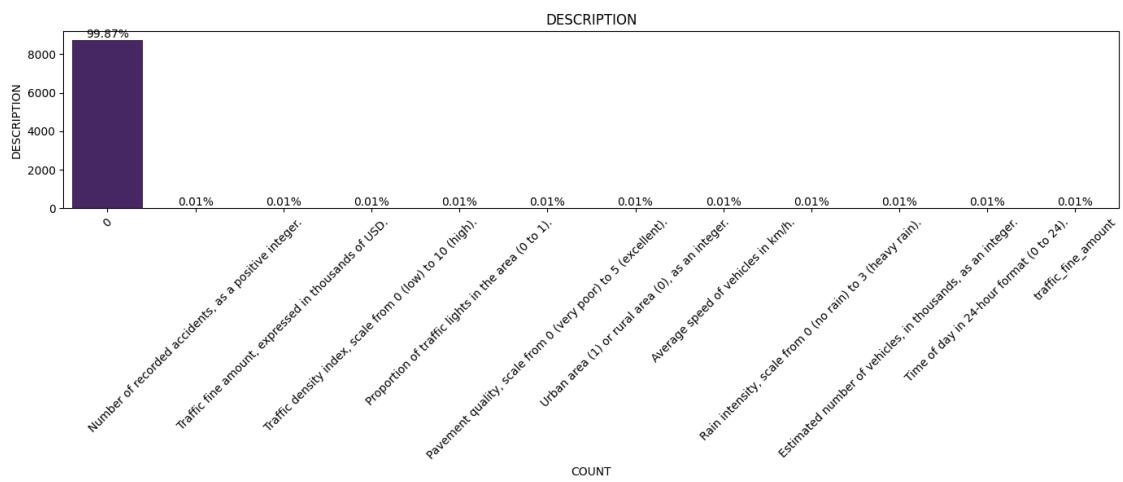
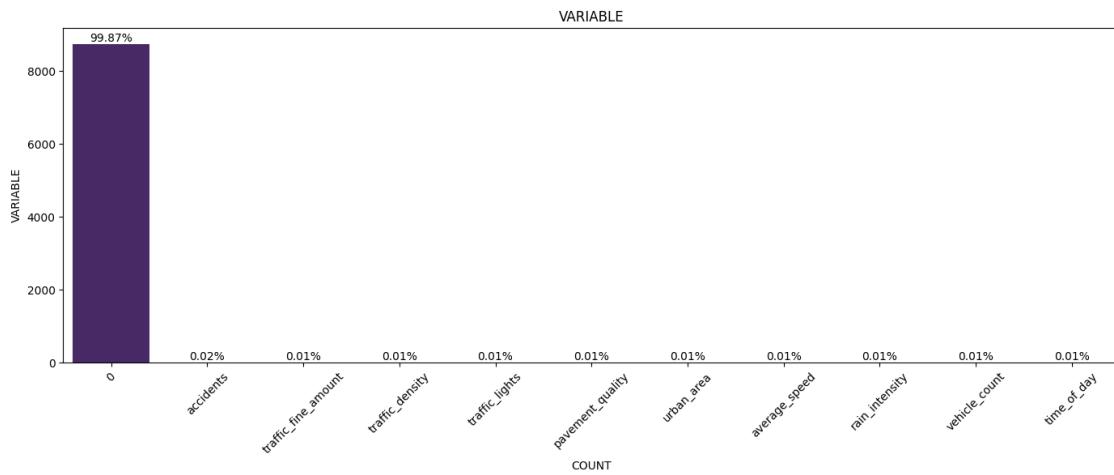
```
ACCIDENTS      False
TRAFFIC_FINE_AMOUNT  False
```

```

TRAFFIC_DENSITY    False
TRAFFIC_LIGHTS     False
PAVEMENT_QUALITY   False
URBAN_AREA         False
AVERAGE_SPEED      False
RAIN_INTENSITY     False
VEHICLE_COUNT      False
TIME_OF_DAY        False

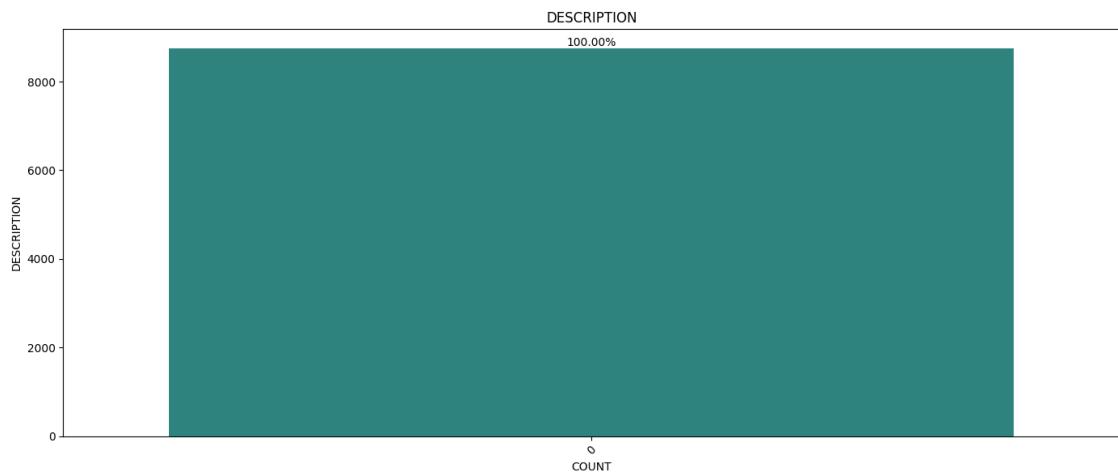
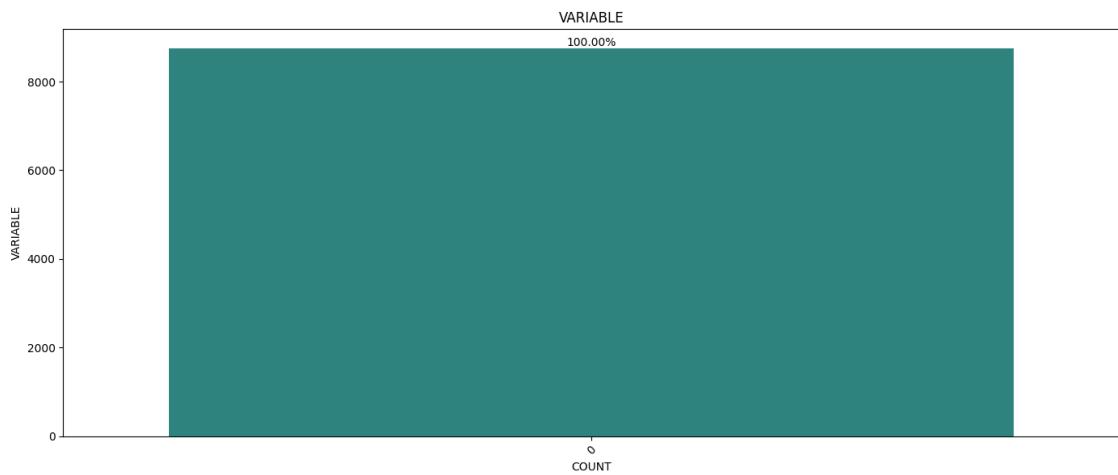
```

```
[24]: for col in cat_cols:
    cat_plot(df,col)
```

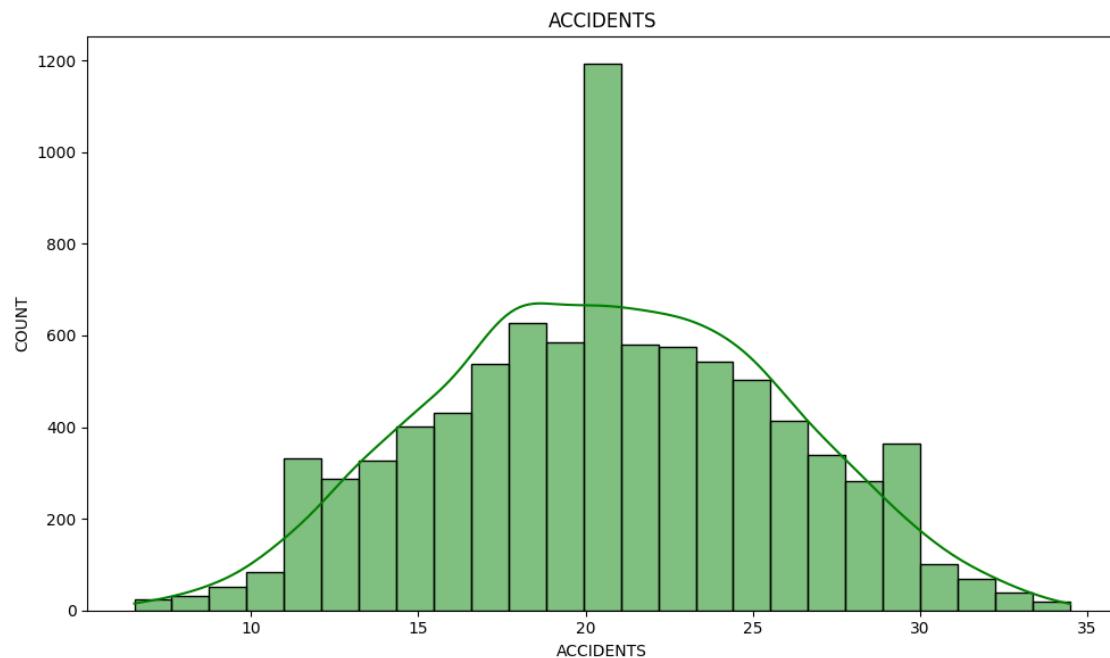


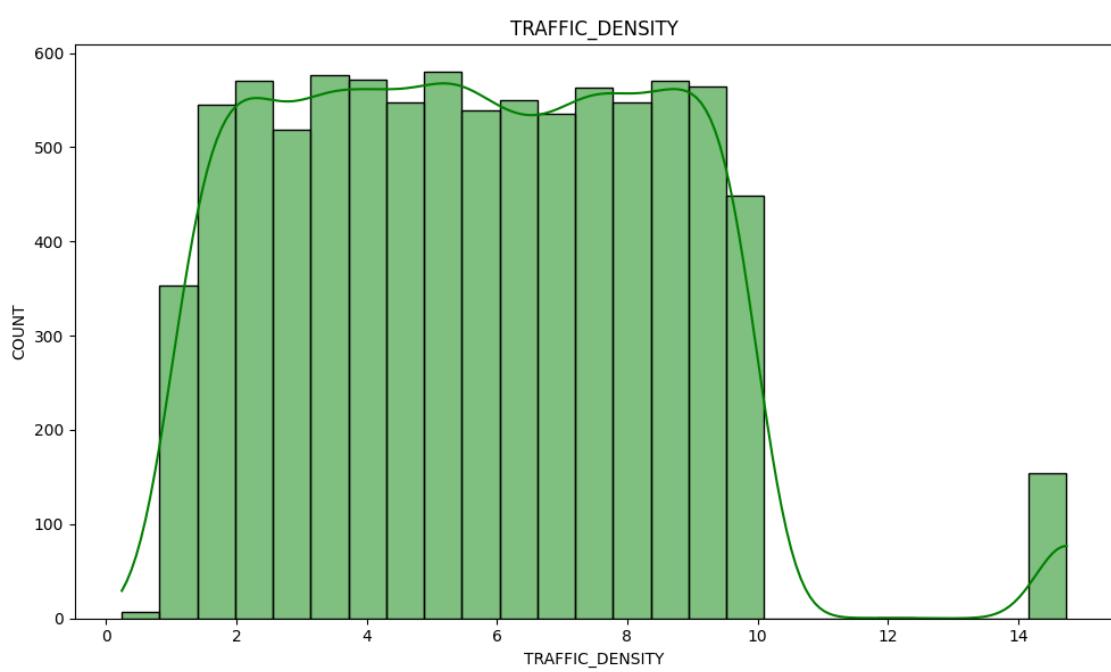
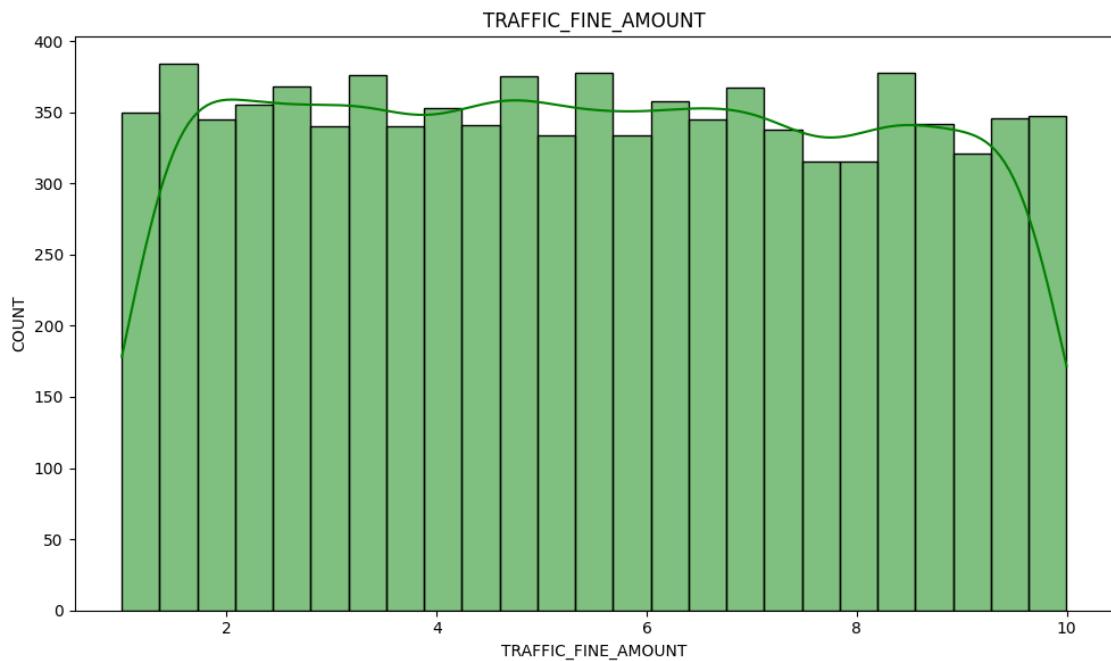
```
[25]: df = df[df["Variable"]=="0"]
```

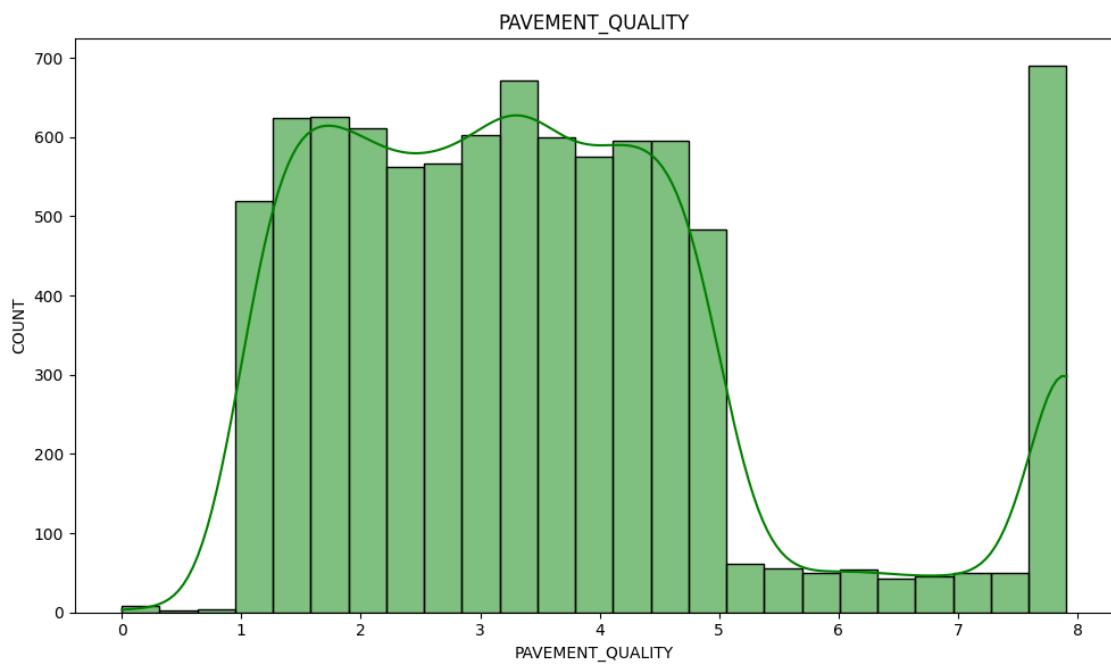
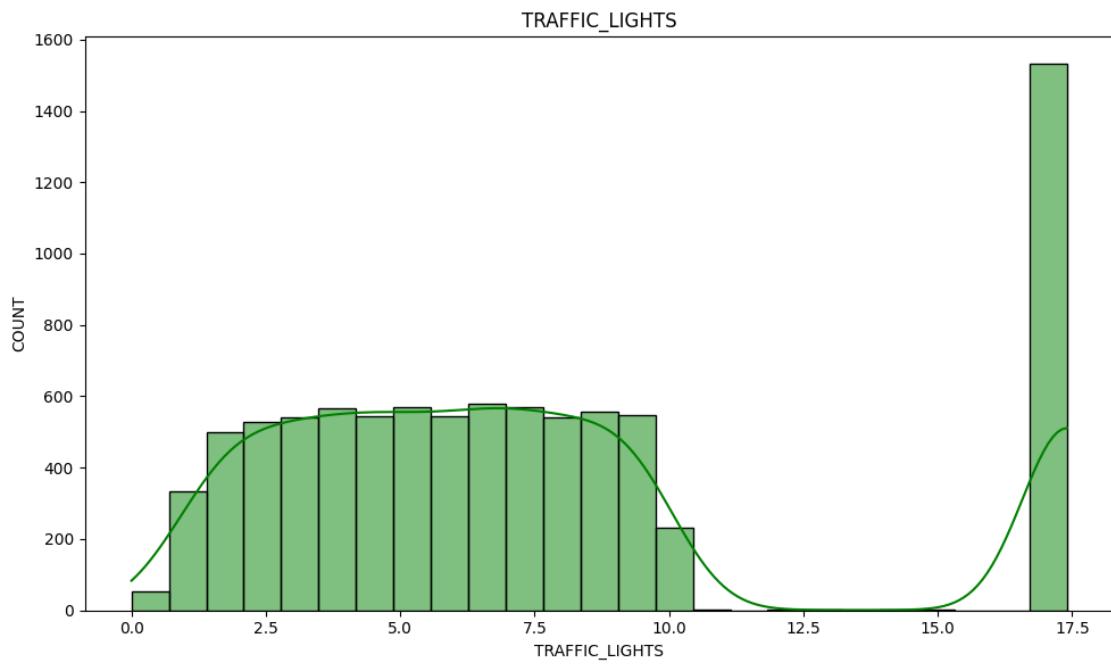
```
[26]: for col in cat_cols:  
    cat_plot(df,col)
```

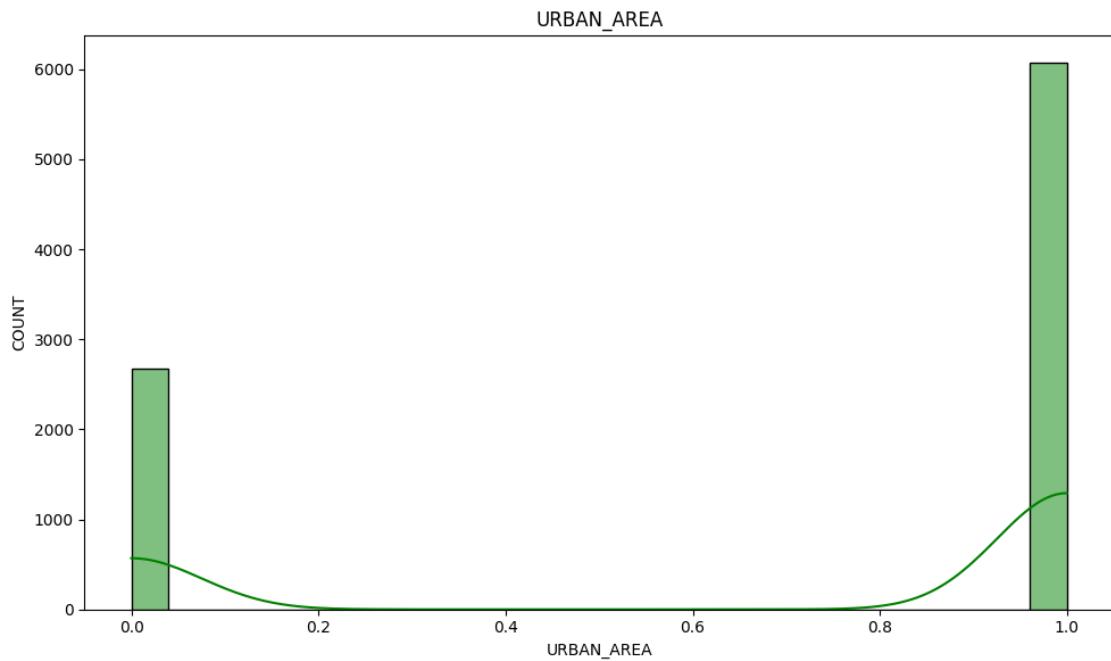


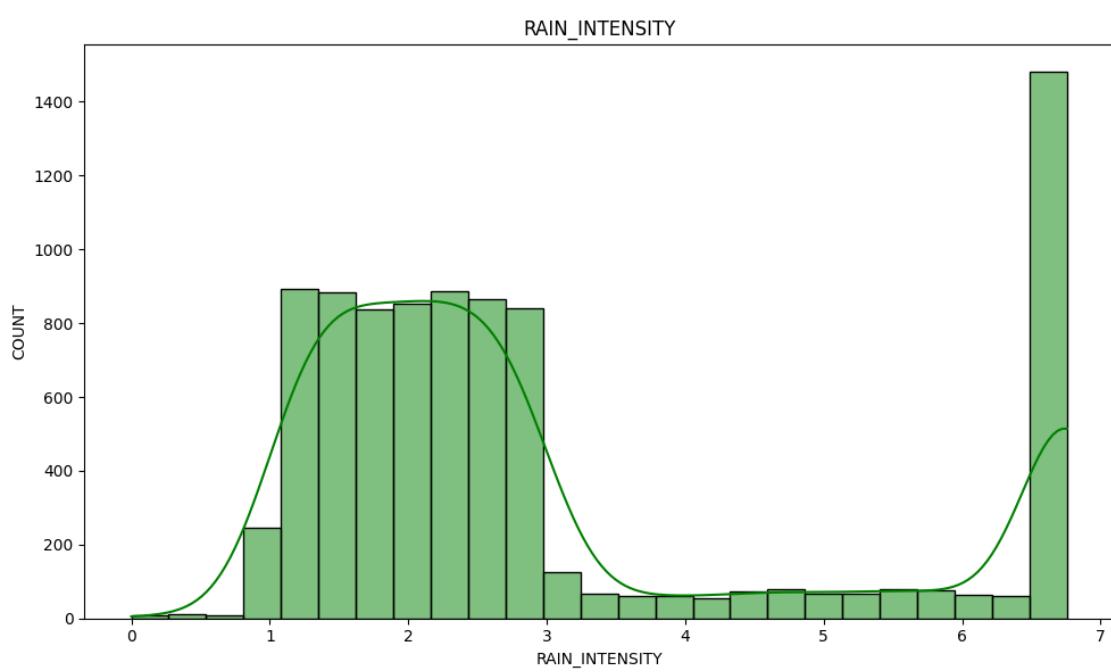
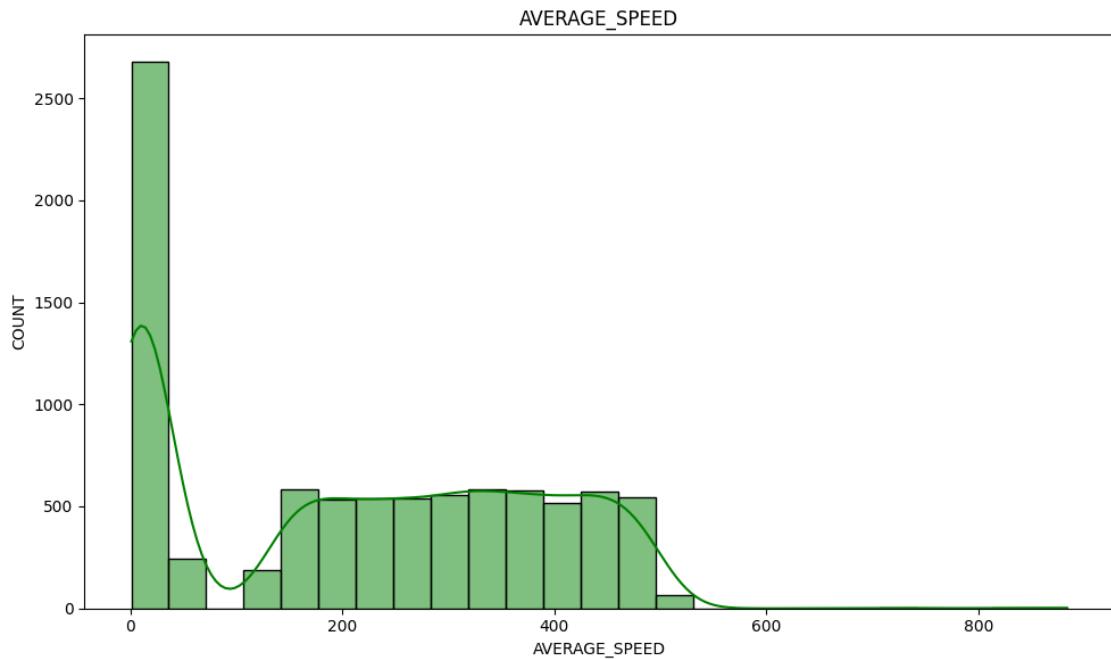
```
[27]: for col in num_cols:  
    hist_plot(df,col)
```

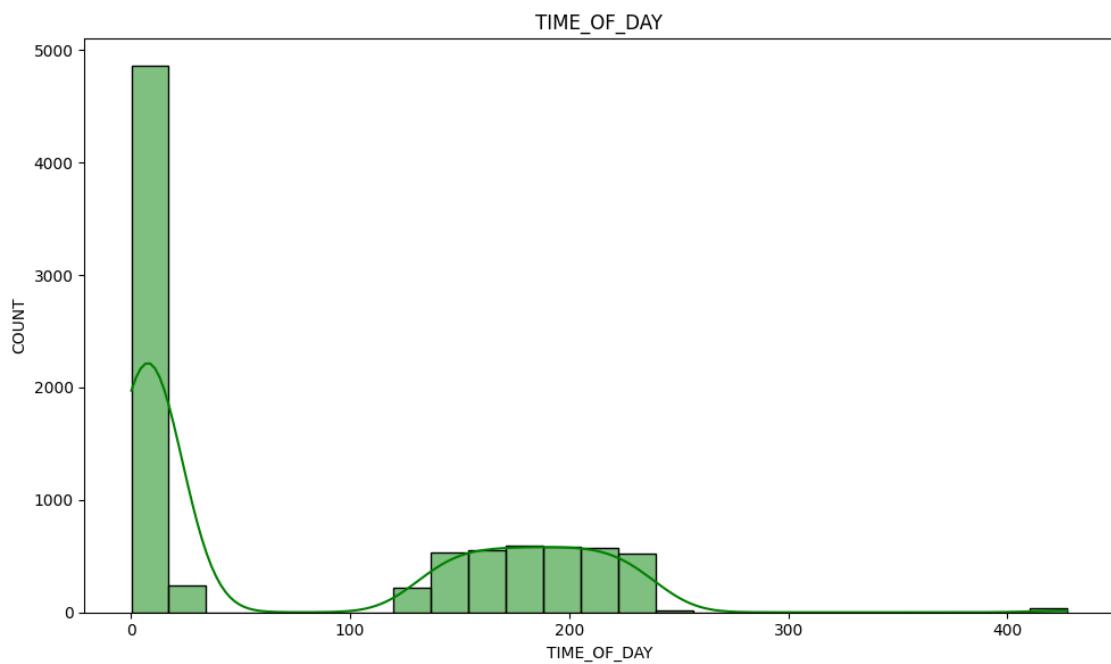
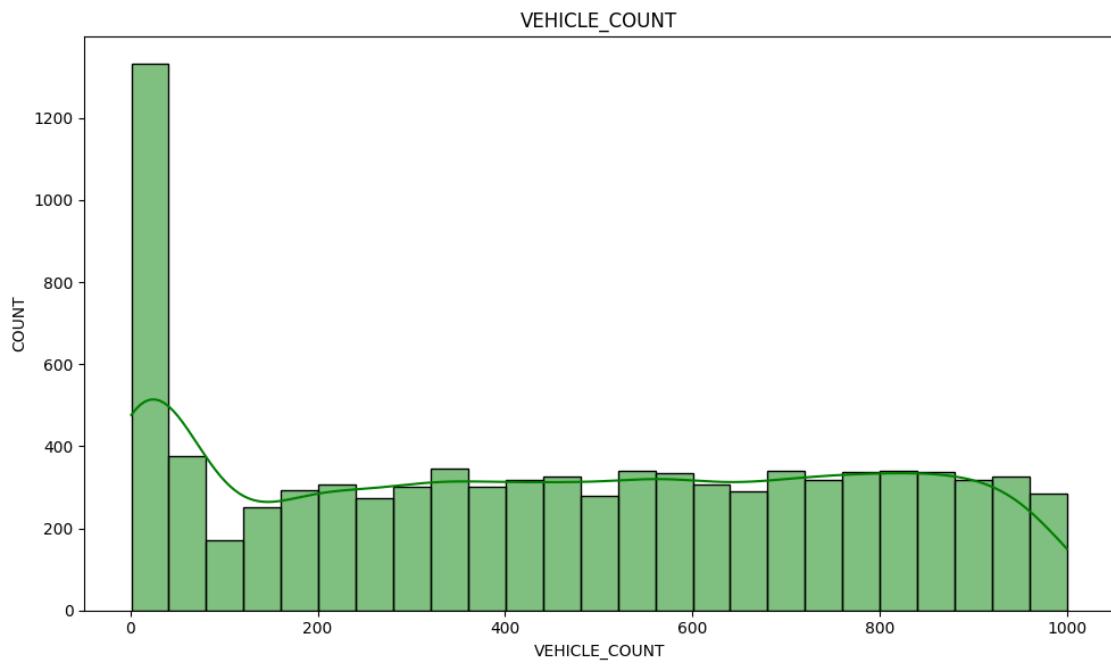




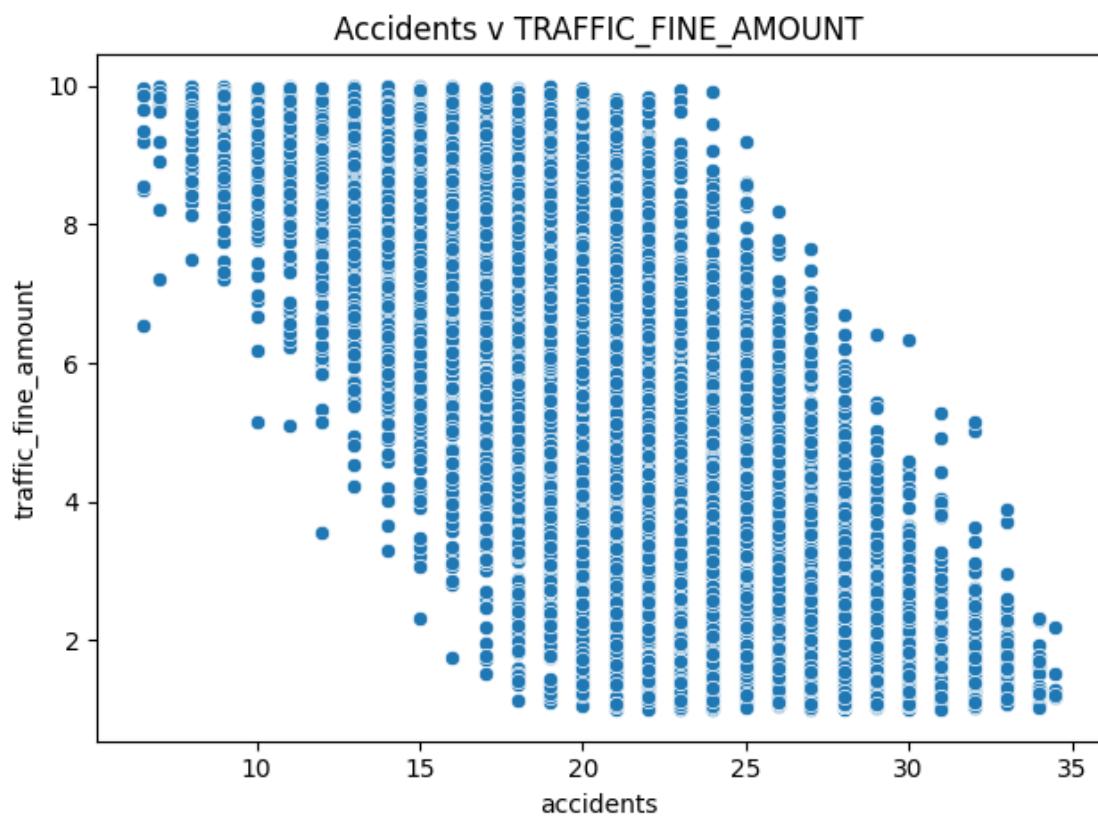


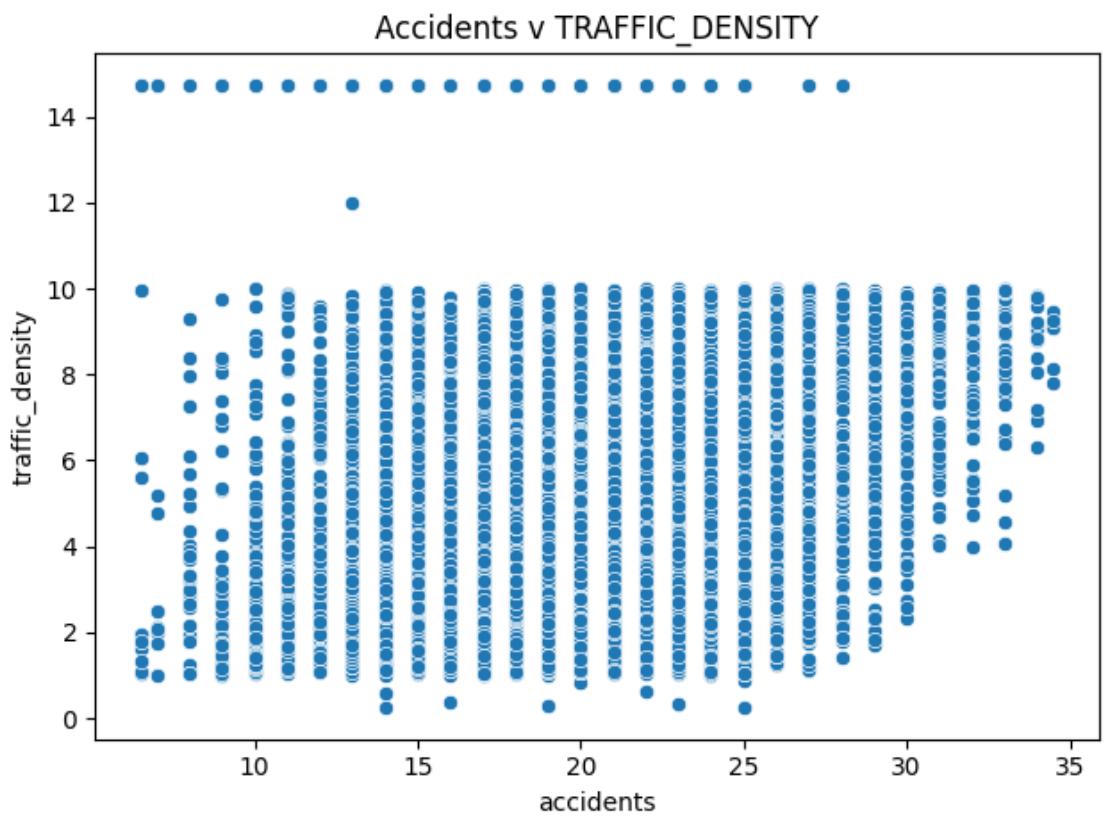


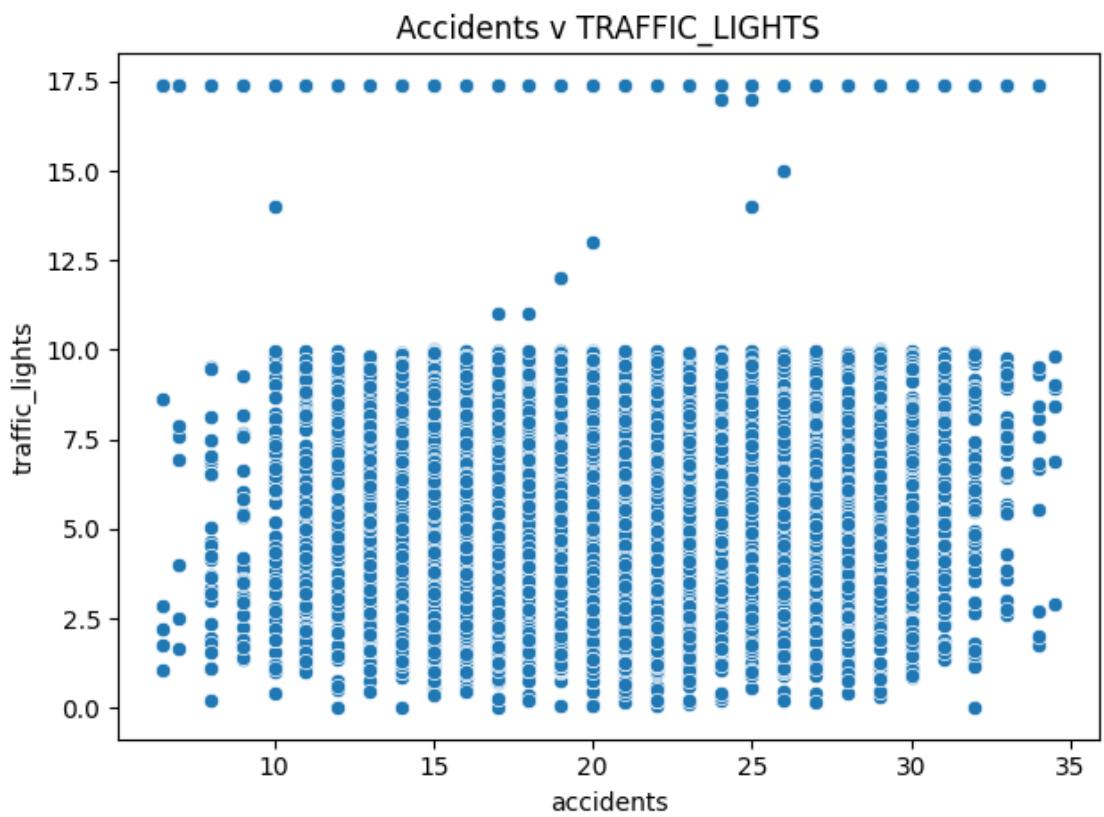




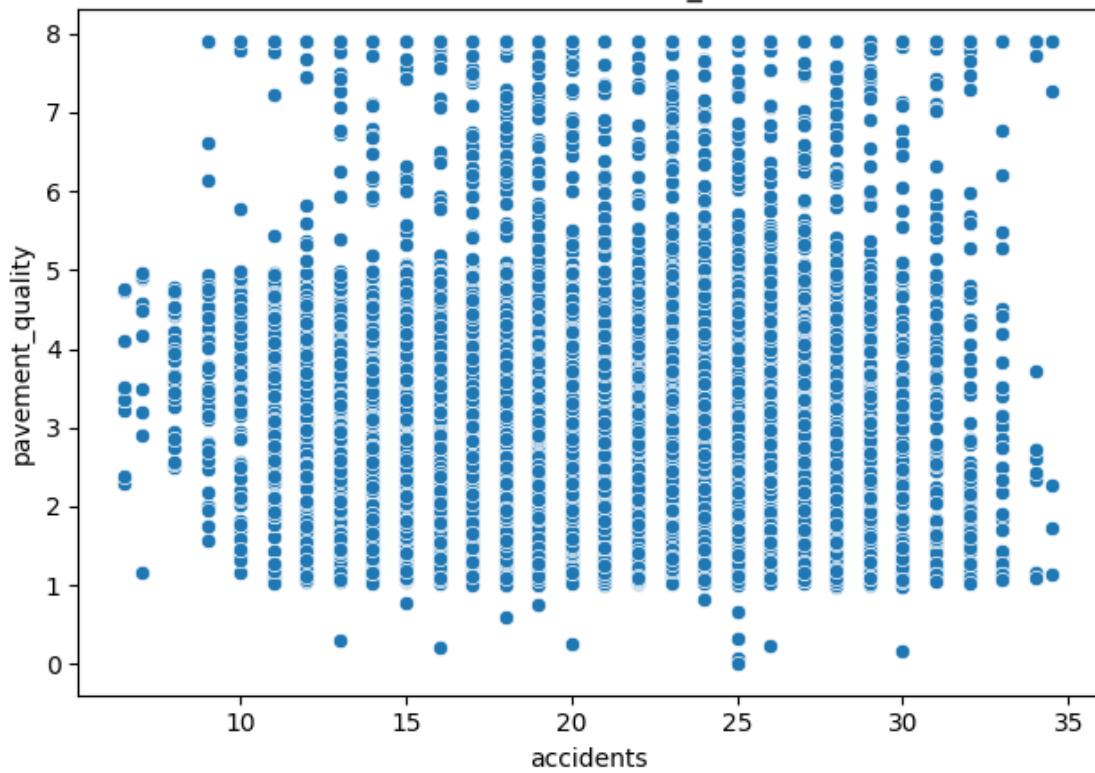
```
[28]: nums = [col for col in num_cols if col not in "accidents"]
for col in nums:
    sns.scatterplot(data=df, x="accidents", y=col)
    plt.title(f"Accidents v {col.upper()}")
    plt.tight_layout()
    plt.show()
```

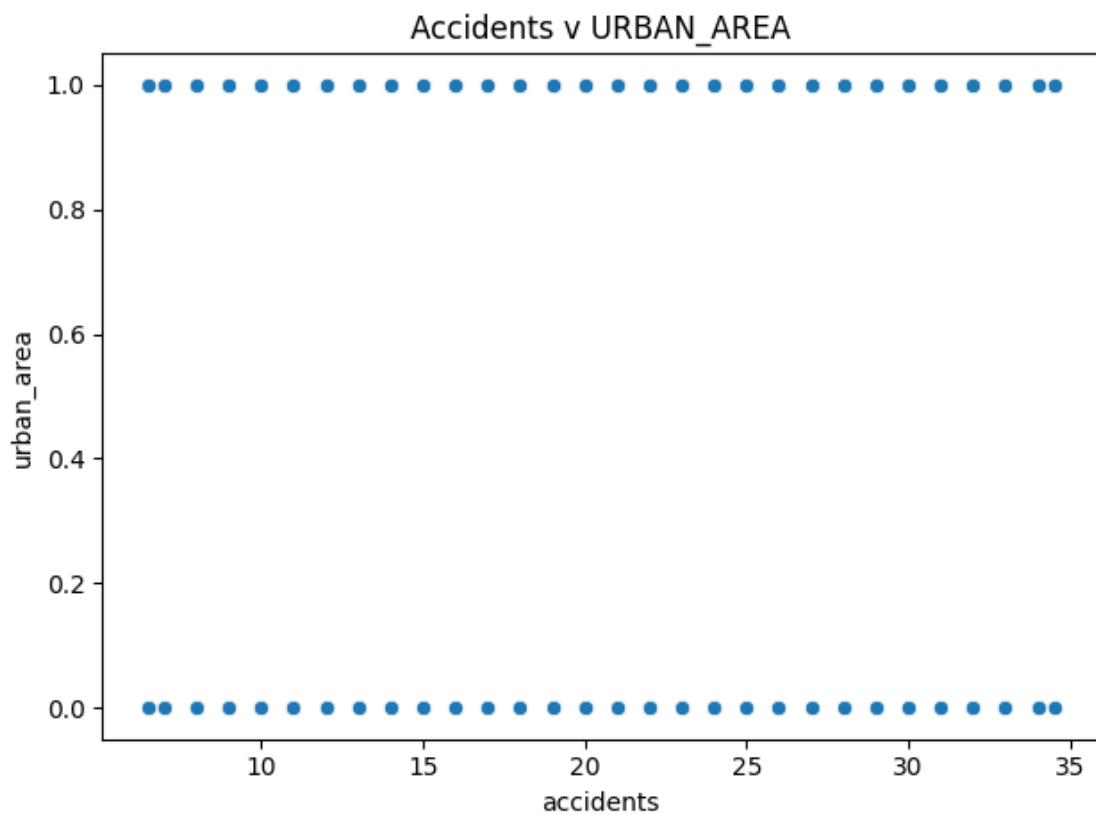


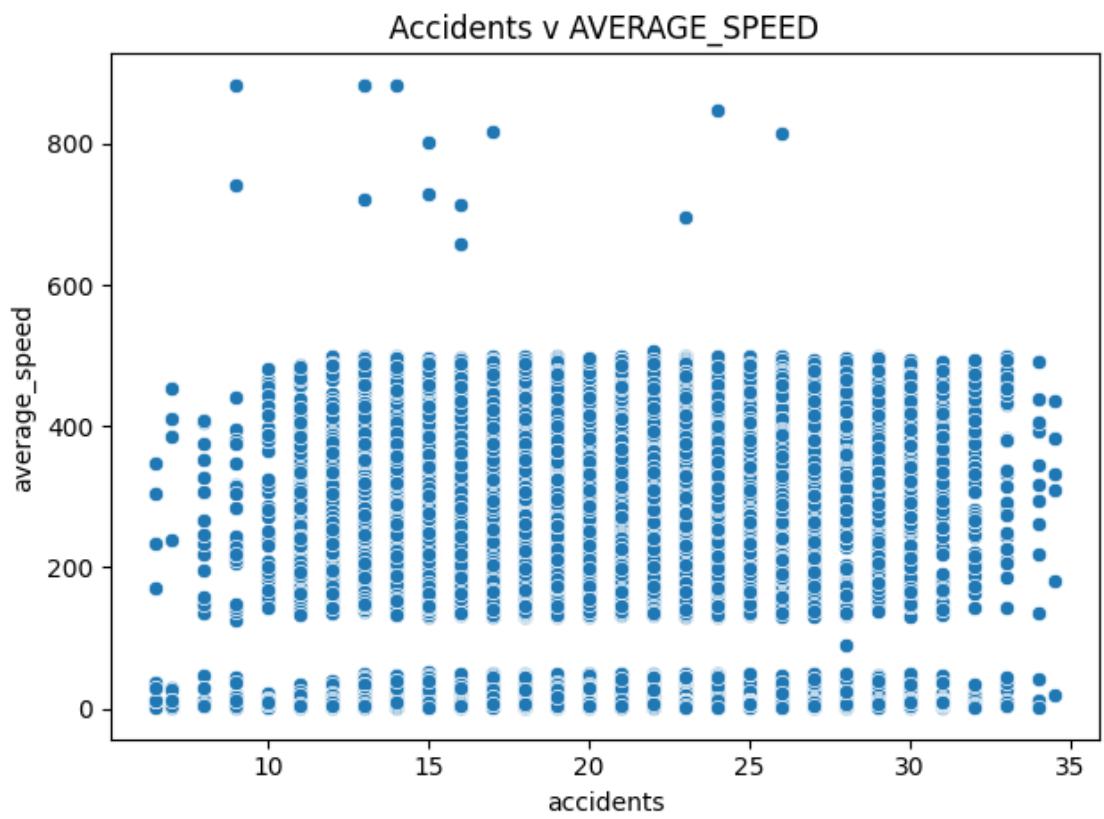




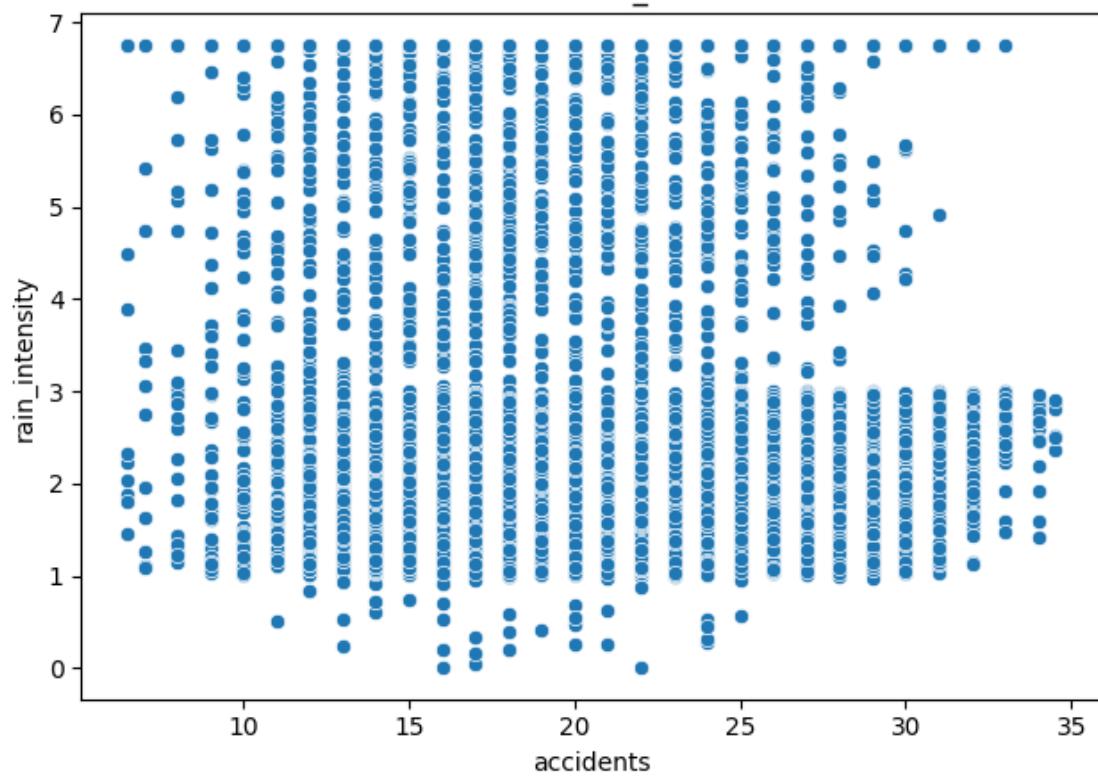
Accidents v PAVEMENT_QUALITY

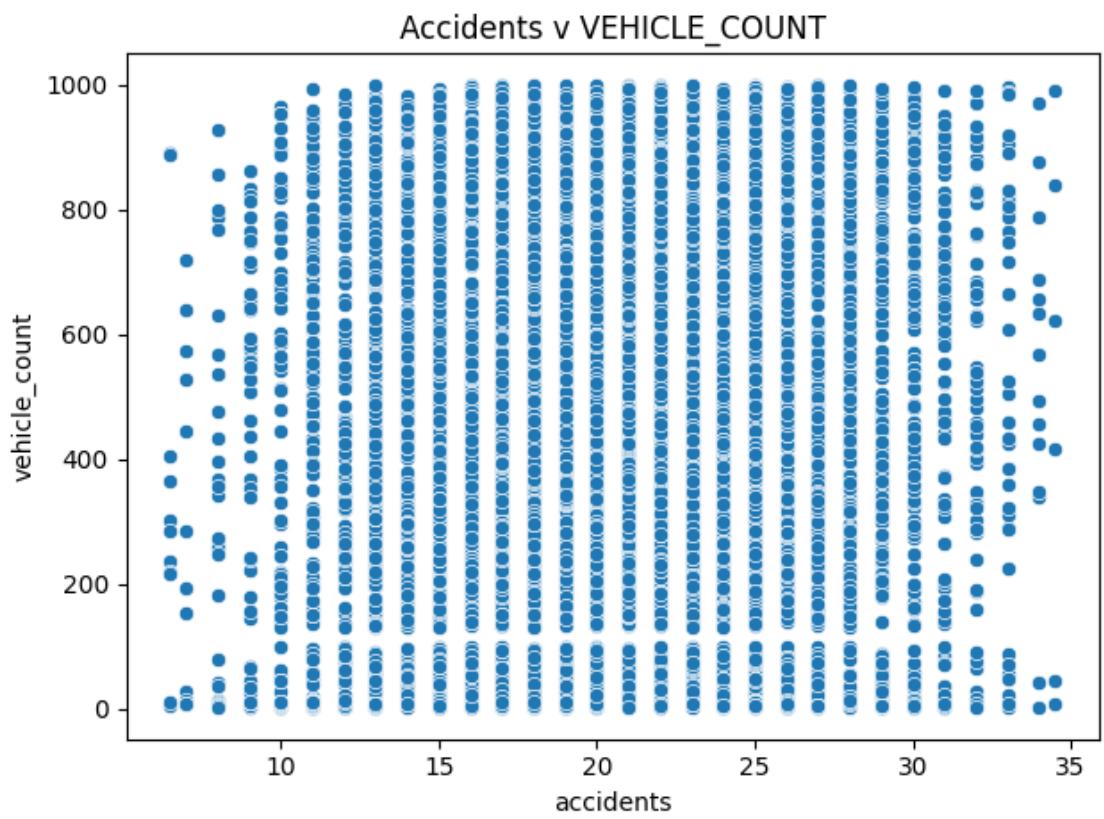


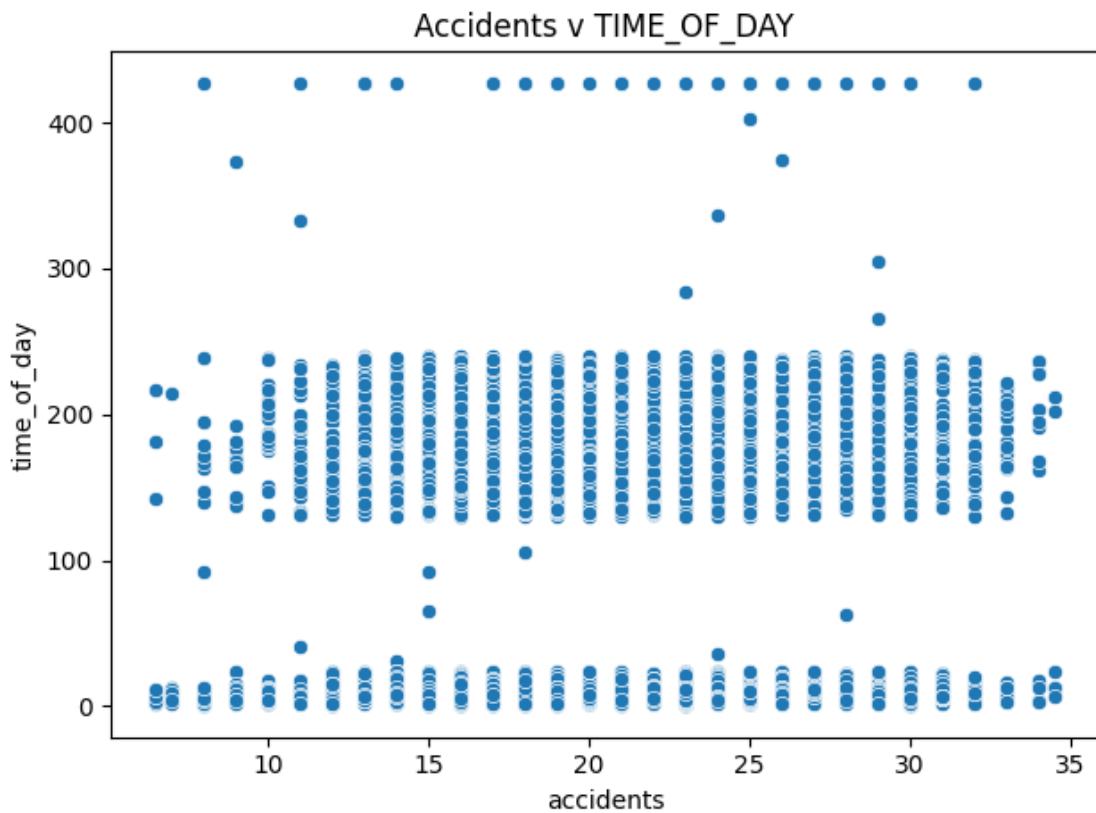




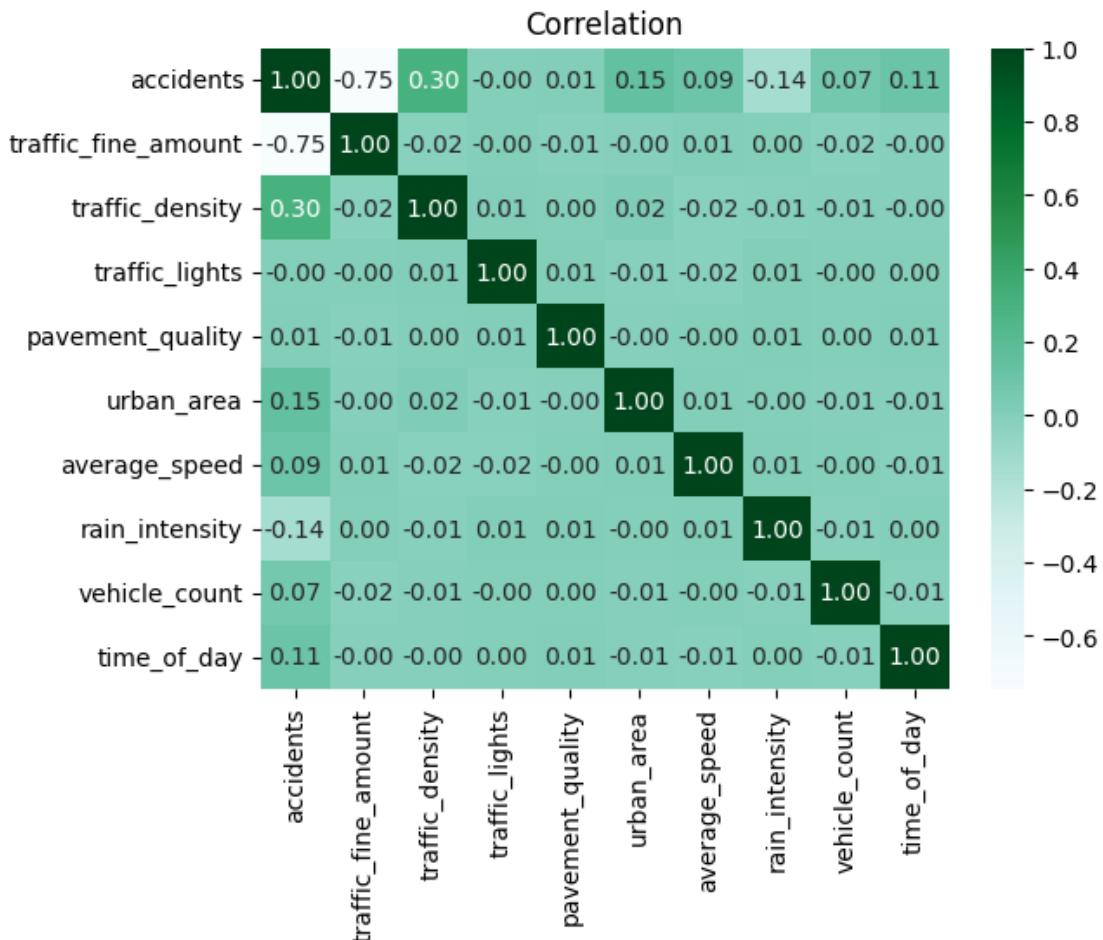
Accidents v RAIN_INTENSITY







```
[29]: sns.heatmap(df[num_cols].corr(), annot=True, fmt=".2f", cmap="BuGn")
plt.title("Correlation", fontsize=12)
plt.show()
```



```
[30]: X = df.drop(["accidents","Variable","Description"],axis=1)
y = df["accidents"]

[31]: X_train , X_test , y_train , y_test = train_test_split(X,y,test_size=0.
           ↪20,random_state=42)

[32]: scaler = MinMaxScaler()

[33]: X_train = scaler.fit_transform(X_train)
X_test = scaler.fit_transform(X_test)

[34]: print(f"X Train Shape : {X_train.shape}\nX Test Shape : {X_test.shape}\nY Train_
           ↪Shape : {y_train.shape}\nY Test Shape : {y_test.shape}")
```

X Train Shape : (6996, 9)
X Test Shape : (1749, 9)
Y Train Shape : (6996,)
Y Test Shape : (1749,)

```
[35]: base_models(X_train,y_train,"r2")
```

```
Base Models
r2: 0.6958 Ridge
r2: 0.116 Lasso
r2: 0.0833 ElasticNet
r2: 0.6958 Linear
r2: 0.5707 Decision Tree
r2: 0.801 Random Forest
r2: 0.8352 Gradient Boosting
r2: 0.808 XGB
r2: 0.7134 KNN
r2: 0.7973 MLP
r2: 0.6958 Bayesian Ridge
r2: 0.8466 CatBoost
```

```
[36]: cat_params = {
    'iterations': [100, 200, 300, 500],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'depth': [3, 5, 7, 10]
}

gb_params = {
    'n_estimators': [100, 200, 300, 500],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'max_depth': [3, 5, 7, 10]
}
xgb_params = {
    'n_estimators': [100, 200, 300, 500],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'max_depth': [3, 5, 7, 10]
}
```

```
[37]: best_models = hyperparameter_optimization_randomized(X_train,y_train,n_iter=10)
```

```
HYPERPARAMETER OPTIMIZATION WITH RANDOMIZEDSEARCH
----- GradientBoosting -----
r2 (Before): 0.8361
r2 (After): 0.8389
GradientBoosting best params: {'n_estimators': 100, 'max_depth': 5,
'learning_rate': 0.1}

----- CatBoost -----
r2 (Before): 0.8457
r2 (After): 0.8488
CatBoost best params: {'learning_rate': 0.05, 'iterations': 300, 'depth': 5}

----- XGBoost -----
```

```
r2 (Before): 0.8068  
r2 (After): 0.8309  
XGBoost best params: {'n_estimators': 100, 'max_depth': 5, 'learning_rate': 0.1}
```

```
[38]: voting_reg = voting_regressor(best_models,X_train,y_train)
```

Voting Regressor
Neg MSE: 4.272705632654722
R²: 0.8442041169906338

```
[39]: predict = voting_reg.predict(X_test)
```

```
[40]: print(f'R2 Score: {r2_score(y_test, predict):.4f}')  
print(f'MSE: {mean_squared_error(y_test, predict):.4f}')  
print(f'RMSE: {np.sqrt(mean_squared_error(y_test, predict)):.4f}')
```

R2 Score: 0.8453
MSE: 4.1274
RMSE: 2.0316

```
[41]: plot_importance_ensemble(voting_reg,X)
```

